

Overview of the `DataStorage` Class

****Package:**** `com.data_management`

****Purpose:****

The `DataStorage` class serves as the central repository for managing and retrieving patient data in the cardiovascular data simulator. It organizes patient records by patient ID, storing `Patient` objects in a `HashMap` and providing methods to add data, retrieve records within a time range, and access all patients. The class acts as a bridge between data generation (from `com.cardio_generator.generators`) and data processing (e.g., alert generation in `com.alerts`), enabling structured storage and access to health metrics like heart rate, blood pressure, and ECG.

****Role in the Project:****

The `DataStorage` class is the backbone of the simulator's data management subsystem, coordinating the storage of data produced by `PatientDataGenerator` implementations (e.g., `ECGDataGenerator`, `BloodPressureDataGenerator`) and facilitating data retrieval for analysis or monitoring. It interacts with the `Patient` and `PatientRecord` classes to manage patient-specific records and supports Task 7's alert system by providing data to `com.alerts.AlertGenerator` for evaluation, which may trigger `Alert` objects via `AlertFactory` and `AlertDecorator`. The class is used by `HealthDataSimulator` to store data output via `OutputStrategy` implementations (e.g., `FileOutputStrategy`, `WebSocketOutputStrategy`) and supports both real-time and historical data access, critical for clinical monitoring and alert generation.

****Technical Characteristics:****

- The class uses a `HashMap` to store `Patient` objects indexed by `patientId`, enabling $O(1)$ access to patient data.
- It provides methods for adding data (`addPatientData`), retrieving records (`getRecords`), and accessing all patients (`getAllPatients`), supporting CRUD-like operations.
- The class includes a `main` method for demonstration, though it references an undefined `DataReader`, indicating integration with external data sources.

- It is designed for concurrent use in `HealthDataSimulator`'s multi-threaded environment but lacks explicit thread safety, requiring careful handling in concurrent scenarios.

Internal Components and Their Purposes

The `DataStorage` class consists of a field, a constructor, three data management methods, and a `main` method. Each component is described below, including its purpose, technical details, and role in the class's functionality.

1. **Field:**

- **private Map<Integer, Patient> patientMap**
 - **Purpose:** Stores `Patient` objects indexed by their unique `patientId`, serving as the primary data structure for organizing patient data.
 - **Technical Details:**
 - Type: `Map<Integer, Patient>`, implemented as a `HashMap` for O(1) lookup, insertion, and deletion by `patientId`.
 - Private access ensures encapsulation, initialized in the constructor and modified via `addPatientData` and accessed via `getRecords` and `getAllPatients`.
 - Keys are `Integer` (boxed `int` patient IDs), and values are `Patient` objects containing patient-specific records.
 - **Role:** Provides efficient storage and retrieval of patient data, enabling the system to manage multiple patients and their records.

2. **Constructor:**

- **public DataStorage()**
 - **Purpose:** Initializes a new `DataStorage` instance with an empty `HashMap` for storing patient data.

- **Technical Details:**
 - No parameters, as the class starts with an empty state.
 - Logic: Initializes `patientMap` as a new `HashMap<Integer, Patient>`.
 - Public access allows instantiation by `HealthDataSimulator` or other components (e.g., the `main` method).
 - **Role:** Prepares the `DataStorage` instance for use, setting up the underlying storage structure to accept patient data.

3. **Methods:**

- **`public void addPatientData(int patientId, double measurementValue, String recordType, long timestamp)`**
 - **Purpose:** Adds or updates patient data by creating or retrieving a `Patient` object and adding a new record to its record list.
 - **Technical Details:**
 - Parameters:
 - `patientId` (`int`), the unique identifier of the patient.
 - `measurementValue` (`double`), the numerical value of the health metric (e.g., 120.0 for blood pressure).
 - `recordType` (`String`), the type of measurement (e.g., "HeartRate", "BloodPressure").
 - `timestamp` (`long`), the time of measurement in milliseconds since the Unix epoch.
 - Return Type: `void`, as the method's effect is updating `patientMap` and the patient's records.
 - Logic:
 - Retrieves the `Patient` for `patientId` from `patientMap` using `get`.
 - If no `Patient` exists (`null`), creates a new `Patient` with `patientId` and adds it to `patientMap` using `put`.
 - Calls `patient.addRecord` to add a new `PatientRecord` with the provided `measurementValue`, `recordType`, and `timestamp`.

- Public access allows data generators or data readers (e.g., ``DataReader``) to store data via ``DataStorage``.
 - Does not validate inputs (e.g., positive ``patientId``, non-null ``recordType``), assuming valid data from callers.
 - **Role:** Enables the storage of health data, creating or updating patient records as data is generated or ingested, forming the basis for data management.
-
- **``public List<PatientRecord> getRecords(int patientId, long startTime, long endTime)``**
 - **Purpose:** Retrieves a list of ``PatientRecord`` objects for a specific patient, filtered by a time range.
 - **Technical Details:**
 - Parameters:
 - ``patientId`` (``int``), the unique identifier of the patient.
 - ``startTime`` (``long``), the start of the time range in milliseconds since the Unix epoch.
 - ``endTime`` (``long``), the end of the time range in milliseconds since the Unix epoch.
 - Return Type: ``List<PatientRecord>``, a list of records within ``[startTime, endTime]`` or an empty list if the patient is not found.
 - Logic:
 - Retrieves the ``Patient`` for ``patientId`` from ``patientMap`` using ``get``.
 - If the ``Patient`` exists, calls ``patient.getRecords(startTime, endTime)`` to get filtered records.
 - If no ``Patient`` is found, returns an empty ``ArrayList``.
 - Public access allows ``com.alerts.AlertGenerator``, monitoring tools, or other components to retrieve records for analysis or alert generation.
 - Does not validate ``patientId``, ``startTime``, or ``endTime`` (e.g., ensuring ``startTime <= endTime``), relying on ``Patient.getRecords`` for filtering logic.
 - **Role:** Facilitates time-based data retrieval for a specific patient, supporting analysis of health trends or alert generation based on historical data.

- **``public List<Patient> getAllPatients()``**
 - **Purpose:** Retrieves a list of all `Patient` objects stored in the `DataSource`.
 - **Technical Details:**
 - No parameters, as it returns all patients.
 - Return Type: `List<Patient>`, a new `ArrayList` containing all `Patient` objects from `patientMap.values()`.
 - Logic: Creates a new `ArrayList` from `patientMap.values()` to avoid exposing the internal `HashMap`'s collection.
 - Public access allows `com.alerts.AlertGenerator` or monitoring tools to iterate over all patients for system-wide analysis or alert evaluation.
 - **Role:** Enables access to all patient data, supporting bulk operations like evaluating all patients for alerts or generating summary statistics.

- **``public static void main(String[] args)``**
 - **Purpose:** Demonstrates the usage of `DataSource` by initializing the system, simulating data ingestion, retrieving records, and triggering alert evaluations.
 - **Technical Details:**
 - Parameters: `args` (`String[]`), command-line arguments (unused in the provided code).
 - Return Type: `void`, as it is a demonstration method.
 - Logic:
 - Comments indicate an undefined `DataReader` that should read data into `storage` (e.g., from a file or network source), but it is not implemented.
 - Creates a new `DataSource` instance.
 - Retrieves records for `patientId` 1 within a hardcoded time range (1700000000000L to 1800000000000L) and prints them to `System.out`, showing `patientId`, `recordType`, `measurementValue`, and `timestamp`.
 - Creates an `AlertGenerator` instance with the `DataSource`.
 - Iterates over all patients (via `getAllPatients()`) and calls `alertGenerator.evaluateData` to check for alert conditions.

- Static and public access allows execution as a standalone program for testing or demonstration.
- References an undefined ``DataReader`` and ``PatientRecord`` methods (e.g., ``getPatientId``, ``getRecordType``), indicating dependencies on other classes.
- **Role:** Provides a proof-of-concept for ``DataStorage`` usage, demonstrating data storage, retrieval, and integration with the alert system, though it requires additional components (``DataReader``) to be fully functional.

Technical Points and Design Considerations

1. **Data Organization and Efficiency:**

- The use of a ``HashMap<Integer, Patient>`` ensures $O(1)$ access to ``Patient`` objects by ``patientId``, making ``addPatientData`` and ``getRecords`` efficient for patient lookup. The ``ArrayList`` in ``Patient`` for records complements this by providing $O(1)$ appends and $O(n)$ retrieval, suitable for typical simulation workloads.
- The separation of concerns between ``DataStorage`` (managing patients) and ``Patient`` (managing records) promotes modularity, though the redundant ``patientId`` in ``PatientRecord`` could be optimized by relying on the ``Patient`` context.

2. **Thread Safety:**

- The ``patientMap`` (``HashMap``) is not thread-safe, and concurrent calls to ``addPatientData``, ``getRecords``, or ``getAllPatients`` in ``HealthDataSimulator``'s multi-threaded environment could cause race conditions (e.g., ``ConcurrentModificationException`` or inconsistent ``Patient`` creation).
- To ensure thread safety, ``patientMap`` could be replaced with a ``ConcurrentHashMap``, or methods could be synchronized, though this may introduce performance overhead. Alternatively, ``DataStorage`` could be designed as a singleton or use a dedicated writer thread to serialize updates.

3. **Error Handling and Validation:**

- The class lacks input validation for `patientId`, `measurementValue`, `recordType`, `timestamp` (`addPatientData`), and `startTime`/`endTime` (`getRecords`). For example, negative `patientId`, null `recordType`, or invalid time ranges could lead to inconsistent data. Adding validation (e.g., throwing `IllegalArgumentException`) would improve robustness.
- The `getRecords` method gracefully handles missing patients by returning an empty list, but it could log such cases for debugging. The `main` method assumes `PatientRecord` methods exist, indicating incomplete integration.

4. **Performance Considerations:**

- The `addPatientData` method is efficient ($O(1)$ for `HashMap` operations, $O(1)$ for `Patient.addRecord`), but `getRecords` depends on `Patient.getRecords`' $O(n)$ iteration, which could be slow for large record sets. Indexing records by timestamp (e.g., in a `TreeMap`) or using a database could improve retrieval performance.
- The `getAllPatients` method creates a new `ArrayList`, which is $O(n)$ for n patients but avoids exposing `patientMap`'s internal collection, ensuring encapsulation.
- For large-scale simulations, integrating with a database (e.g., SQLite, PostgreSQL) via JDBC could enhance scalability and persistence.

5. **Integration with Task 7 Design Patterns:**

- The `DataStorage` class is central to Task 7's alert system, providing data to `com.alerts.AlertGenerator` via `getRecords` and `getAllPatients`. The `evaluateData` method in `AlertGenerator` processes `Patient` records to trigger alerts, which are created by `AlertFactory` subclasses (e.g., `ECGAlertFactory`) and enhanced by `AlertDecorator` subclasses (e.g., `PriorityAlertDecorator`).
- The generic `recordType` in `addPatientData` supports diverse data types, including alerts (e.g., `recordType` "Alert"), aligning with Task 7's flexible alert processing.
- The class integrates with the Strategy Pattern via `OutputStrategy` implementations (e.g., `WebSocketOutputStrategy`), which may output records retrieved by `getRecords` for real-time monitoring or logging, supporting alert dissemination.

6. **Potential Improvements:**

- Add input validation for `patientId` (positive), `recordType` (non-null), `timestamp` (non-negative), and time ranges (`startTime <= endTime`) to throw `IllegalArgumentException` for invalid values.
- Make `patientMap` thread-safe (e.g., using `ConcurrentHashMap`) to support concurrent access in multi-threaded environments.
- Optimize `getRecords` for large datasets by indexing records in `Patient` or integrating with a database for faster queries.
- Implement a `removePatient` or `pruneRecords` method to manage memory by deleting old data or unused patients in long-running simulations.
- Replace `System.out` in the `main` method with a logging framework (e.g., SLF4J) for structured logging and better debugging.
- Define an interface for `DataStorage` (e.g., `DataStore`) to support alternative implementations (e.g., database-backed storage) via dependency injection.
- Complete the `main` method by implementing or mocking `DataReader` and ensuring `PatientRecord` method compatibility (e.g., `getPatientId`, `getRecordType`).

Interaction with Other Components

- **With `Patient` Class:** The `DataStorage` class manages a collection of `Patient` objects, creating them in `addPatientData` and delegating record storage (`addRecord`) and retrieval (`getRecords`) to `Patient` instances. The `getAllPatients` method returns `Patient` objects for system-wide operations.
- **With `PatientRecord` Class:** `DataStorage` indirectly interacts with `PatientRecord` via `Patient`, as `addPatientData` creates records through `Patient.addRecord`, and `getRecords` retrieves `PatientRecord` objects. The `main` method assumes `PatientRecord` methods like `getPatientId` and `getRecordType`.
- **With `PatientDataGenerator` Implementations:** Data generators (e.g., `ECGDataGenerator`, `BloodPressureDataGenerator`) in `com.cardio_generator.generators`

produce data that `HealthDataSimulator` stores in `DataStorage` via `addPatientData`, mapping generator output to `measurementValue`, `recordType`, and `timestamp`.

- **With `com.alerts.AlertGenerator` (Task 7):** The `DataStorage` class provides data to `AlertGenerator` via `getRecords` (for specific patients) and `getAllPatients` (for system-wide evaluation). `AlertGenerator` processes records to trigger alerts, which are created by `AlertFactory` subclasses and enhanced by `AlertDecorator` subclasses.

- **With `Alert` Class (Task 7):** Alerts generated from `DataStorage` records (e.g., stored as `recordType` "Alert") are represented as `Alert` objects (from `com.alerts`), supporting Task 7's Factory and Decorator patterns.

- **With `OutputStrategy` Implementations:** Records retrieved by `getRecords` can be output via `OutputStrategy` implementations (e.g., `WebSocketOutputStrategy`, `FileOutputStrategy`) for real-time monitoring or logging, supporting Task 7's alert dissemination.

- **With `DataReader` (Undefined):** The `main` method references an undefined `DataReader` class, likely responsible for reading external data (e.g., from files or networks) into `DataStorage` via `addPatientData`. This indicates integration with external data sources, though it requires implementation.

Summary of Functionality

The `DataStorage` class is the central hub of the cardiovascular data simulator's data management subsystem, organizing patient data in a `HashMap` of `Patient` objects and providing methods to add data (`addPatientData`), retrieve records (`getRecords`), and access all patients (`getAllPatients`). It integrates with `HealthDataSimulator` to store data from `PatientDataGenerator` implementations and supports Task 7's alert system by supplying data to `com.alerts.AlertGenerator` for evaluation, compatible with `AlertFactory` and `AlertDecorator`. The class's use of `HashMap` ensures efficient patient access, and its `main` method demonstrates usage, though it requires a `DataReader` implementation. The class is effective for managing clinical data, but enhancements in thread safety, validation, and performance optimization could improve its scalability for large-scale simulations. Overall, `DataStorage` plays a pivotal role in enabling structured data management, supporting real-time monitoring, historical analysis, and alert generation in the simulator.

